

Tipps for coding exercise on PPO


Prof. Dr. Maximilian Schiffer

Professorship of Operations and Supply Chain Management

TUM School of Management

Technical University of Munich

WiSe21/22



3rd coding exercise: PPO

- Find optimal pole balancing policy using proximal policy optimization (PPO) → same problem as for the coding exercise presented this week
- Use the same gym environment as the one that you used for the exercise from last week, so you can compare your results
- Try to think about the different ingredients that you will need:
 - A neural network for the actor and for the critic
 - A memory/buffer where you store states, actions, rewards from one trajectory through an episode & that for N actors
 - A function that calculates advantage estimates
 - A function that updates weights of both critic and actor
- You can use the code framework introduced this week or implement according to your own taste

Pseudo code for PPO-Clip

Initialize network weights w, θ_{old}

Repeat:

For N actors:

$$\tau \leftarrow \{s_1, a_1, \dots, s_T, a_T, r_T, s_{T+1}\} \sim \pi(s, a; \theta_{\text{old}})$$

Compute advantage estimates A_1, \dots, A_T based on $V(s; w)$


$$\theta_{\text{new}} \leftarrow \arg \max_{\theta_{\text{new}}} \frac{1}{NT} \sum_{\tau} \sum_t L_t^{\text{CLIP}}$$

Update w

$$\theta_{\text{old}} \leftarrow \theta_{\text{new}}$$

n-step temporal-difference learning

Idea: we assume a tabular setting for now and let TD target look n steps into the future




New update rule

- We define the n-step return $G_t^{(n)} = \sum_{k=t}^{t+n-1} \gamma^{k-t} \cdot r_k + \gamma^n V(s_{t+n})$
- We obtain a new update rule for the tabular value function estimate: $V(s_t) \leftarrow V(s_t) + \alpha (G_t^{(n)} - V(s_t))$

Forward-view TD(λ): averaging n-step returns

Idea: we combine all n-step returns $G_t^{(n)}$ into one weighted average, the λ -return G_t^λ



New update rule

- We define the λ -return
$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$
- For episodes with terminal state T


$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_t^{(n)} + \lambda^{T-t-1} G_t$$
- We obtain a new update rule called forward-view TD(λ):
$$V(s_t) \leftarrow V(s_t) + \alpha (G_t^\lambda - V(s_t))$$

PPO-Clip: the effect of clipping

Case 1: positive advantage

- The objective term reduces to

$$L^{\text{CLIP}} = \min \left(\frac{\pi_{\theta_{\text{new}}}(s, a)}{\pi_{\theta_{\text{old}}}(s, a)}, 1 + \epsilon \right) \cdot A_{\pi_{\theta_{\text{old}}}}(s, a)$$




- Since the advantage is positive, maximizing the objective means to maximize $\pi_{\theta_{\text{new}}}$
- Due to the min, increasing the ratio of old and new policy r beyond $1 + \epsilon$ does not change the objective

[10]

Case 2: negative advantage

- The objective term reduces to

$$L^{\text{CLIP}} = \max \left(\frac{\pi_{\theta_{\text{new}}}(s, a)}{\pi_{\theta_{\text{old}}}(s, a)}, 1 - \epsilon \right) \cdot A_{\pi_{\theta_{\text{old}}}}(s, a)$$



- Since the advantage is negative, maximizing the objective means to minimize $\pi_{\theta_{\text{new}}}$
- Due to the max, decreasing the ratio of old and new policy r below $1 - \epsilon$ does not change the objective

[10]

Moving the new policy too far from the old one does not have a positive effect on the objective